

```
In [ ]: # Import needed libraries

import findspark
findspark.init('/usr/hdp/2.6.5.0-292/spark2')

# Create a Spark Context which will be used for distributed data

import pyspark
sc = pyspark.SparkContext(appName="Twitter Topic Sentiment")

import string

import re as re

import nltk

import time

from pyspark.sql import SQLContext

from pyspark.sql.types import *

from pyspark.sql.functions import monotonically_increasing_id

from pyspark.mllib.util import MLUtils

from pyspark.ml.feature import RegexTokenizer, Tokenizer, StopWordsRemover

from pyspark.mllib.clustering import LDA, LDAModel

nltk.download('stopwords')

from nltk.corpus import stopwords

from pyspark.mllib.linalg import Vector as oldVector, Vectors as oldVectors

from pyspark.ml.linalg import Vector as newVector, Vectors as newVectors

from pyspark.ml.feature import IDF

import numpy as np

import matplotlib.pyplot as plt

import pyspark.sql.functions as func
```

```
In [ ]: # Create an SQL Context which will be used for sql like distribu
        # As I get more familiar with what technology to use where I wil
        # pyspark dataframes, and pandas dataframes
```

```
sqlContext = SQLContext(sc)
```

```
In [ ]: # Hadoop is the filesystem being used. This is a three node virt
        # Read in data from Hadoop
```

```
ITData = sqlContext.readFile("hdfs://user/username/topicsum/output")
```

```
In [ ]: # Output sample of data
```

```
ITData.take(5)
```

```
In [ ]: # Count number of records loaded to pyspark RDD
```

```
ITData.count()
```

```
In [ ]: # By default, data is partitioned based on the data size
```

```
# Check the number of partitions created
```

```
ITData.getNumPartitions()
```

```
In [ ]: # Twitter data was collected and batched in files with each file
```

```
# Extract the first file header from the dataset and display
```

```
# This will be used later to remove all headers from the dataset
```

```
header = ITData.first()
```

```
header
```

```
In [ ]: # Filter all of the headers from the data set
```

```
# Count the number of records remaining in the data set
```

```
# If 10 files were read from Hadoop, this count should be 10 less
```

```
ITData_NoHeader = ITData.filter(lambda row : row != header)
```

```
ITData_NoHeader.count()
```

```
In [ ]: # We now have an RDD with not header information

# In preparation for creating a dataframe from the RDD, create a
schema = StructType([
    StructField('timetext', StringType(), nullable=True),
    StructField('tweet_id', StringType(), nullable=True),
    StructField('tweet_source', StringType(), nullable=True),
    StructField('tweet_truncated', StringType(), nullable=True),
    StructField('tweet_text', StringType(), nullable=True),
    StructField('tweet_user_screen_name', StringType(), nullable=True),
    StructField('tweet_user_id', StringType(), nullable=True),
    StructField('tweet_user_location', StringType(), nullable=True),
    StructField('tweet_user_description', StringType(), nullable=True),
    StructField('tweet_user_followers_count', StringType(), nullable=True),
    StructField('tweet_user_statuses_count', StringType(), nullable=True),
    StructField('tweet_user_time_zone', StringType(), nullable=True),
    StructField('tweet_user_geo_enabled', StringType(), nullable=True),
    StructField('tweet_user_lang', StringType(), nullable=True),
    StructField('tweet_coordinates_coordinates', StringType(), nullable=True),
    StructField('tweet_place_country', StringType(), nullable=True),
    StructField('tweet_place_country_code', StringType(), nullable=True),
    StructField('tweet_place_full_name', StringType(), nullable=True),
    StructField('tweet_place_name', StringType(), nullable=True),
    StructField('tweet_place_type', StringType(), nullable=True)
])

# Create a dataframe from the RDD with schema

ITData_df = sqlContext.createDataFrame(ITData_NoHeader.map(lambda
ITData_df.printSchema()
```

```
In [ ]: # First convert dataframe to rdd

# Use map lambda to select the tweet_text column and filter out
tweet = ITData_df.rdd.map(lambda document: document['tweet_text']).filter(lambda
```

```
In [ ]: # Retrieve stop words. Note we may need to add to the stop words
StopWords = stopwords.words("english")
```

```
In [ ]: # Further clean tweets, split them out into individual words, and

tokens = tweet.map(lambda document: document.strip().lower()) \
    .map(lambda document: re.split(" ", document)) \
    .map(lambda word: [x for x in word if x.isalpha()]) \
    .map(lambda word: [x for x in word if len(x) > 3]) \
    .map(lambda word: [x for x in word if x not in StopWords])
```

```
In [ ]: # tokens is an RDD, display the first 5 records
```

```
tokens.take(5)
```

```
In [ ]: # Create a new dataframe from the above RDD, adding column names
```

```
tweet_df = sqlContext.createDataFrame(tokens, ["tweet_words"])
```

```
In [ ]: # Display the first 5 records of the dataframe
```

```
tweet_df.show(5)
```

```
In [ ]: # Prepare for Topic Modeling
```

```
print(time.strftime('%m%d%Y %H:%M:%S'))
cv = CountVectorizer(inputCol="tweet_words", outputCol="raw_features")
cvmodel = cv.fit(tweet_df)
print(time.strftime('%m%d%Y %H:%M:%S'))
```

```
In [ ]: print(time.strftime('%m%d%Y %H:%M:%S'))
```

```
result_cv = cvmodel.transform(tweet_df)
```

```
print(time.strftime('%m%d%Y %H:%M:%S'))
```

```
In [ ]: result_cv.show(1)
```

```
In [ ]: rs = result_cv.rdd.map(lambda (x, y): (x, y.toDenseVector.fromM
```

```
In [ ]: rs_df = rs.toDF(["tweet_index", "index", "raw_features"])
```

```
In [ ]: rs.take(1)
```

```
In [ ]: rs_df.show(1)
```

```
In [ ]: print(time.strftime('%m%d%Y %H:%M:%S'))
```

```
idf = IDF(inputCol="raw_features", outputCol="features")
```

```
idfModel = idf.fit(result_cv)
```

```
result_tfidf = idfModel.transform(result_cv)
```

```
print(time.strftime('%m%d%Y %H:%M:%S'))
```

```
In [ ]: # Run the LDA Topic Modeler
```

```
# Note the time before and after is printed in order to find out
```

```
print(time.strftime('%m%d%Y %H:%M:%S'))
```

```
num_topics = 10
```

```
max_iterations = 20
```

```
lda_model = LDA.train(rs_df['index', 'raw_features'].rdd.map(lambda
```

```
print(time.strftime('%m%d%Y %H:%M:%S'))
```

```
In [ ]: vocabArray = cvmodel.vocabulary
```

```
In [ ]: # Set the top number of topics to write to spark
```

```
wordNumbers = 20
```

```
In [ ]: def topic_render(topic):
```

```
    terms = topic[0]
```

```
    result = []
```

```
    for i in range(wordNumbers):
```

```
        term = vocabArray[terms[i]]
```

```
        result.append(term)
```

```
    return result
```

```
In [ ]: print(time.strftime('%m%d%Y %H:%M:%S'))
```

```
topics_final = topicIndices.map(lambda topic:
```

```
                                topic_render(topic)).collect()
```

```
print(time.strftime('%m%d%Y %H:%M:%S'))
```

```
In [ ]: # Display topics
```

```
for topic in range(len(topics_final)):
```

```
    print("Topic" + str(topic) + ":")
```

```
    for term in topics_final[topic]:
```

```
        print(term)
```

```
    print("\n")
```

```
In [ ]: # The above above relates topics to the terms I searched in Twit
```

```
# For sentiment analysis, I would like to rate the actual search
```

```
# For this I will build a python array with those search terms
```

```
search_terms = ["machine_learning", "computer_programmer", "data  
"data_scientist", "systems_engineer", "data_anal  
"web_programmer", "automation_engineer", "data_p  
"software_engineer", "software_developer", "info  
"business_intelligence", "enterprise_architect",  
"information_technology", "data", "java", "iot",  
"etl", "devops", "cloud", "developer", "programm
```

```
search_terms
```

```
In [ ]: # Python function to search for topics within a tweet
```

```
# Function will return the topic and the related tweet or NA is
```

```
def SearchTopics(topics, tweet_text):
```

```
    for term in topics:
```

```
        result = tweet_text.find(term)
```

```
        if result > -1:
```

```
            return term, tweet_text
```

```
    return "NA", tweet_text
```

```
In [ ]: # While removing stopwords helps obtain valid topics it will not
        # With topics in hand, topics_final, we will use tweets where st
        tweet.take(5)
```

```
In [ ]: # Search each tweet for topics returning only tweets that match
        # SearchTopics will return both the topic and the related tweet
        # Sentiment will be done on these tweets
        topic_tweet = tweet.map(lambda x: SearchTopics(search_terms=x))
```

```
In [ ]: # Display 5 topic tweet combinations
        topic_tweet.take(10)
```

```
In [ ]: # Setup sentiment analysis
        import nltk
        from nltk.sentiment.vader import SentimentIntensityAnalyzer
        nltk.download('vader_lexicon')
```

```
In [ ]: # Python function to print the sentiment scores
        # This function will have topic and related tweet as in put
        # This function will perform sentiment analysis and output topic
        # Also note this function will only return the compound portion
        # Revert sigpipe to default behavior

        def print_sentiment_scores(topic, sentence):
            snt = SentimentIntensityAnalyzer().polarity_scores(sentence)
            print("{:-<40} {}".format(sentence, str(snt)))
            print(str(snt))
            return(topic, sentence, str(snt.get('compound')))
```

```
In [ ]: # Retrieve sentiment for each topic, tweet
        topic_tweet_sentiment = topic_tweet.map(lambda x: print_sentimen
```

```
In [ ]: # Display sentiment
        topic_tweet_sentiment.take(10)
```

```
In [ ]: # Assign the topic and sentiment only
        topic_tweet_sentiment_pair = topic_tweet_sentiment.map(lambda x:
```

```
In [ ]: # Display topic, sentiment combination
```

```
topic_tweet_sentiment_pair.take(10)
```

```
In [ ]: # Convert to dataframe naming columns
```

```
topic_tweet_sentiment_pair_df = topic_tweet_sentiment_pair.toDF()
```

```
In [ ]: # Display dataframe
```

```
topic_tweet_sentiment_pair_df.show(5)
```

```
In [ ]: # Count sentiment records
```

```
topic_tweet_sentiment_pair_df.count()
```

```
In [ ]: # Create panda dataframe based on topic, sentiment dataframe
# This dataframe will enable us to plot highs, lows, and means
```

```
pdf1 = topic_tweet_sentiment_pair_df.toPandas()
```

```
In [ ]: # Check new dataframe types
```

```
pdf1.dtypes
```

```
In [ ]: # Sentiment is currently of type object, needs to be float
```

```
# Convert sentiment datatype to float
```

```
pdf1['sentiment'] = pdf1.sentiment.astype(float)
```

```
# Check datatypes
```

```
pdf1.dtypes
```

```
# list new panda dataframe
```

```
pdf1
```

```
In [ ]: # Describe data
```

```
pdf1.describe()
```

```
In [ ]: topic_groups = pdf1.groupby('topic')
```

```
topic_groups.size()
```

```
In [ ]: sentiment_terms1 = ['loil', 'dota', 'technology', 'loland']
```

```
In [ ]: pdf2 = pdf1[pdf1.topic.isin(sentiment_terms1)]
```

```
pdf2
```

```
In [ ]: # Histograms
```

```
pdf2 sentiment11.plot(kind='hist', bins=5)
```

```
In [ ]: pdf2.plot -- topic_group.plot(kind='bar')
```

```
In [ ]: # Boxplot sentiments by topic
```

```
pdf2.boxplot(by='topic', column=['sentiment11', 'sentiment10'])
```

```
In [ ]:
```